

## Section 1: Data set description & Objective

### 1.1: Introduction of the dataset

On April 15, 1912, during her maiden voyage, the widely considered “unsinkable” RMS Titanic sank after colliding with an iceberg. Unfortunately, there weren’t enough lifeboats for everyone onboard, resulting in the death of 1502 out of 2224 passengers and crew. While there was some element of luck involved in surviving, it seems some groups of people were more likely to survive than others.

**Source:** <https://www.kaggle.com/c/titanic/data> (<https://www.kaggle.com/c/titanic/data>)

### 1.2: Dataset features

We are given with two separate dataset here.

- train.csv: It is a **labelled dataset** which we need to use to train the models.
- test.csv: It is not labelled and used only for predictions.

Both the dataset have following features.

- PassengerId: Unique ID of the passenger
- Pclass: Ticket class
- Name: Full name of the passenger with salutation
- Sex: Gender
- Age: Age in years
- SibSp: Number of siblings / spouses aboard the Titanic
- Parch: Number of parents / children aboard the Titanic
- Ticket: Ticket number
- Fare: Passenger fare
- Cabin: Cabin number
- Embarked: Port of Embarkation

train.csv has the target variable named *Survived*.

### 1.3: Objective

The objective is to make survival prediction on test.csv. Then, the predicted values can be submitted to Kaggle which will calculate the model accuracy.

### 1.4: Glimpse of the dataset

Not to confuse with train-test splitting, we will import train.csv as labelled and test.csv as unlabelled dataset. We will merge both labelled and unlabelled dataset to run Feature engineering on the whole dataset. Before model training, we will break them down again.

```
In [1]: # Importing essentials
import numpy as np, pandas as pd, matplotlib.pyplot as plt, seaborn as sns
```

```
In [2]: # Setting the stage
sns.set_context('notebook')
sns.set_style('white')
```

```
In [3]: # Suppress warning
import warnings
warnings.filterwarnings('ignore')
```

```
In [4]: # Importing dataset
labelled_df = pd.read_csv('train.csv')
unlabelled_df = pd.read_csv('test.csv')
```

```
In [5]: # Checking the shapes  
labelled_df.shape, unlabelled_df.shape
```

```
Out[5]: ((891, 12), (418, 11))
```

Unlabelled dataset has one less columns as it doesn't have target variable.

```
In [6]: # Checking if Labelled dataset all survival values  
labelled_df.Survived.isna().sum()
```

```
Out[6]: 0
```

Good that labelled dataset has no missing target variable values

```
In [7]: # Merging the datasets  
df = pd.concat([labelled_df, unlabelled_df], ignore_index=True)  
df.shape
```

```
Out[7]: (1309, 12)
```

Now, the merged dataset should have 418 missing values in Survived columns.

```
In [8]: df.Survived.isna().sum()
```

```
Out[8]: 418
```

Looking into some samples from the dataset

```
In [9]: df.sample(10)
```

```
Out[9]:
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare
515	516	0.0	1	Walker, Mr. William Anderson	male	47.0	0	0	36967	34.0208
443	444	1.0	2	Reynaldo, Ms. Encarnacion	female	28.0	0	0	230434	13.0000
1223	1224	NaN	3	Thomas, Mr. Tannous	male	NaN	0	0	2684	7.2250
136	137	1.0	1	Newsom, Miss. Helen Monypeny	female	19.0	0	2	11752	26.2833
641	642	1.0	1	Sagesser, Mlle. Emma	female	24.0	0	0	PC 17477	69.3000
1046	1047	NaN	3	Duquemin, Mr. Joseph	male	24.0	0	0	S.O./P.P. 752	7.5500
1243	1244	NaN	2	Dibden, Mr. William	male	18.0	0	0	S.O.C. 14879	73.5000
45	46	0.0	3	Rogers, Mr. William John	male	NaN	0	0	S.C./A.4. 23567	8.0500
570	571	1.0	2	Harris, Mr. George	male	62.0	0	0	S.W./PP 752	10.5000
190	191	1.0	2	Pinsky, Mrs. (Rosa)	female	32.0	0	0	234604	13.0000

## Section 2: Data Analysis

### 2.1: Shape and size

We have missing values for Age, Cabin, Fare and Embarked features. As we have only 891 labelled observations, we cannot afford to remove missing values. We need to find ways to deal with them.

```
In [10]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1309 entries, 0 to 1308
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId     1309 non-null   int64
1   Survived        891 non-null   float64
2   Pclass          1309 non-null   int64
3   Name            1309 non-null   object
4   Sex             1309 non-null   object
5   Age            1046 non-null   float64
6   SibSp           1309 non-null   int64
7   Parch           1309 non-null   int64
8   Ticket          1309 non-null   object
9   Fare            1308 non-null   float64
10  Cabin           295 non-null    object
11  Embarked        1307 non-null   object
dtypes: float64(3), int64(4), object(5)
memory usage: 122.8+ KB
```

We will now go through each of the features except PassengerID to perform feature engineering. We will remove PassengerID as it doesn't add any value to our analysis or modelling.


## 2.2: Pclass

We have no missing data in Pclass. First class passengers are more likely to survive followed by second and third class.

```
In [11]: df.Pclass.value_counts()
```

```
Out[11]: 3    709
         1    323
         2    277
         Name: Pclass, dtype: int64
```

```
In [12]: # Bar plot to check survival probabilities of each ticket class
ax = sns.barplot(x='Pclass', y='Survived', data=df)
```



## 2.3: Name

Name needn't to have any correlation with survival. However, we will extract salutations from names which will help us to determine missing ages.

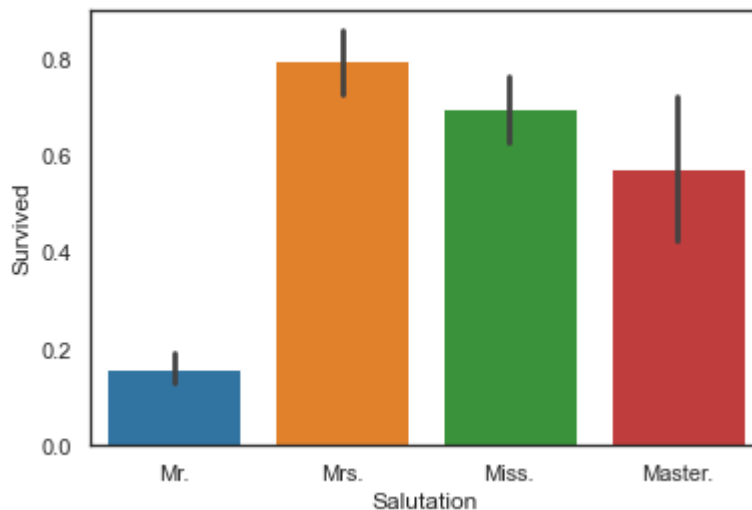
```
In [13]: df['Salutation'] = df['Name'].map(lambda name: name.split(',')[1].split(' ')[1])
df.groupby(['Salutation', 'Sex']).size()
```

```
Out[13]: Salutation Sex
Capt.      male      1
Col.        male      4
Don.        male      1
Dona.       female    1
Dr.         female    1
           male      7
Jonkheer.   male      1
Lady.       female    1
Major.      male      2
Master.     male     61
Miss.       female   260
Mlle.       female     2
Mme.        female     1
Mr.         male    757
Mrs.        female   197
Ms.         female     2
Rev.        male      8
Sir.        male      1
the         female     1
dtype: int64
```

```
In [14]: # Replace salutations for unmarried women with Miss
df['Salutation'].replace('Mlle.', 'Miss.', inplace=True)
# Replace salutations for married women with Mrs
df['Salutation'].replace(['Dona.', 'Lady.', 'Ms.', 'Mme.', 'the'], 'Mrs.', inplace=True)
# Replace female doctor with Mrs
df.loc[(df.Salutation == 'Dr.') & (df.Sex == 'female'), 'Salutation'] = 'Mrs.'
# Replace other salutations with Mr. as they are men
salts = df['Salutation'].value_counts()
male_salts = list(salts[salts < 10].index)
df['Salutation'].replace(male_salts, 'Mr.', inplace=True)
# Final count
df['Salutation'].value_counts()
```

```
Out[14]: Mr.      782
Miss.    262
Mrs.     204
Master.   61
Name: Salutation, dtype: int64
```

```
In [15]: # Bar plot to check survival probabilities of each salutation
ax = sns.barplot(x='Salutation', y='Survived', data=df)
```




Women and kids have very high probability of survival than male. We will further check Sex and Age columns.

## 2.4: Sex

As we have seen above, female are much more likely to survive than male in all classes. However, first class male have more survival rate than others.

```
In [16]: ax = sns.catplot(x='Sex', y='Survived', col='Pclass', kind='bar', data=df)
```



## 2.5: Age


```
In [17]: df.Age.isna().sum()
```

```
Out[17]: 263
```

We need to fill 263 missing values for Ages. Before that, we will see how Age is correlated with Survival probability, just for available Ages.

```
In [18]: fig, ax = plt.subplots(figsize=(10, 5))
ax = sns.kdeplot(data = df.loc[(df.Age.notnull()) & (df.Survived == 0)], 'Age',
                shade = True, color = 'red')
ax = sns.kdeplot(data = df.loc[(df.Age.notnull()) & (df.Survived == 1)], 'Age',
                shade = True, color = 'green')
ax.legend(['Not Survived', 'Survived'])
```

```
Out[18]: <matplotlib.legend.Legend at 0x1de9634c388>
```





The hump between 0 to 10 shows young children are more likely to survive. Also, people aged over 70 have almost no chance to survive.

Now, let fill the missing ages by taking the mean of ages from the group of passengers with same Ticket Class, Sex, Embarked and Salutation.


```
In [19]: df_missing_age = df.loc[df.Age.isna()]

for i, ovn_missing_age in df_missing_age.iterrows():
    pclass = ovn_missing_age.Pclass
    sex = ovn_missing_age.Sex
    embkd = ovn_missing_age.Embarked
    salt = ovn_missing_age.Salutation
    mean_age = df.loc[(df.Pclass == pclass) &
                      (df.Sex == sex) &
                      (df.Embarked == embkd) &
                      (df.Salutation == salt) &
                      (df.Age.notna()), 'Age'].mean()
    df.loc[i, 'Age'] = mean_age
```

Survival density will be impacted after Age approximation, but not hugely. Plotting the graph again to prove.

```
In [20]: fig, ax = plt.subplots(figsize=(10, 5))
ax = sns.kdeplot(data = df.loc[df.Survived == 0, 'Age'], shade = True, color = 'red')
ax = sns.kdeplot(data = df.loc[df.Survived == 1, 'Age'], shade = True, color = 'green')
ax.legend(['Not Survived', 'Survived'])
```

Out[20]: <matplotlib.legend.Legend at 0x1de963caac8>



## 2.6: SibSp


```
In [21]: df.SibSp.value_counts()
```

```
Out[21]: 0      891  
         1      319  
         2       42  
         4       22  
         3       20  
         8        9  
         5         6  
         Name: SibSp, dtype: int64
```

Passengers travelling with small family of one or two members are more likely to survive. Passengers travelling alone also have higher probability to survive than large families.

```
In [22]: sns.catplot(x='SibSp', y='Survived', kind='bar', data=df, aspect=1.5)
```

```
Out[22]: <seaborn.axisgrid.FacetGrid at 0x1de963f9dc8>
```



## 2.7: Parch


```
In [23]: df.Parch.value_counts()
```

```
Out[23]: 0    1002  
         1     170  
         2     113  
         3        8  
         5         6  
         4         6  
         9         2  
         6         2  
         Name: Parch, dtype: int64
```

Passengers travelling with up to 3 children have slightly higher possibility to survive. However, Parch=3 has a very large standard error.

```
In [24]: sns.catplot(x='Parch', y='Survived', kind='bar', data=df, aspect=1.5)
```

```
Out[24]: <seaborn.axisgrid.FacetGrid at 0x1de96481248>
```



## 2.8: Embarked

```
In [25]: df.Embarked.value_counts()
```

```
Out[25]: S    914  
         C    270  
         Q    123  
         Name: Embarked, dtype: int64
```

There are only two missing values in Embarked feature. We will replace those with S, the most frequent value.


```
In [26]: df.Embarked.fillna('S', inplace=True)
df.Embarked.isna().sum()
```

```
Out[26]: 0
```

Passengers from Cherbourg have highest probability of surviving.

```
In [27]: sns.catplot(x='Embarked', y='Survived', kind='bar', data=df, aspect=1.5)
```

```
Out[27]: <seaborn.axisgrid.FacetGrid at 0x1de96139fc8>
```



## 2.9: Ticket

We will categorise all "number only" tickets as Numeric.

```
In [28]: df.loc[df.Ticket.str.isnumeric(), 'Ticket'] = 'NUMERIC'
```

Then, remove . and / from the ticket numbers.

```
In [29]: df['Ticket'] = df['Ticket'].map(lambda tkt: tkt.replace('/', '').replace('.', '').replace(' ', ''))
```

Then, just keep first two characters of the ticket number.

```
In [30]: df['Ticket'] = df['Ticket'].map(lambda tkt: tkt[0:2])
```

Replacing low frequency ticket numbers with Other

```
In [31]: tkt_counts = df['Ticket'].value_counts()
low_tkt_counts = list(tkt_counts[tkt_counts < 20].index)
df['Ticket'].replace(low_tkt_counts, 'OTH', inplace=True)
```


```
In [32]: df['Ticket'].value_counts()
```

```
Out[32]: NU      957
PC       92
CA       69
OTH      68
SO       43
SC       30
A5       28
ST       22
Name: Ticket, dtype: int64
```

PC Ticket holders have highest probability of survival.

```
In [33]: sns.catplot(x='Ticket', y='Survived', kind='bar', data=df, aspect=1.5)
```

```
Out[33]: <seaborn.axisgrid.FacetGrid at 0x1de95e67d88>
```



## 2.10: Fare

We have only one missing fare which we will replace with mean for same Pclass, Sex, Age

```
In [34]: df[df.Fare.isna()]
```

```
Out[34]:
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	E
1043	1044	NaN	3	Storey, Mr. Thomas	male	60.5	0	0	NU	NaN	NaN	

```
In [35]: df.groupby(by=['Pclass', 'Sex', 'Age'])['Fare'].mean()
```

```
Out[35]: Pclass  Sex    Age
1      female  2.0      151.550000
          14.0      120.000000
          15.0      211.337500
          16.0       61.293067
          17.0      82.950000
          ...
3      male    60.5         NaN
          61.0       6.237500
          65.0       7.750000
          70.5       7.750000
          74.0       7.775000
Name: Fare, Length: 369, dtype: float64
```


Closest fare is in between 6 and 7. So, replacing with 6.5

```
In [36]: df.Fare.fillna(6.5, inplace=True)
```

Passengers paying less than 40 ticket fare have very high probability of not surviving.

```
In [37]: fig, ax = plt.subplots(figsize=(10, 5))
ax = sns.kdeplot(data = df.loc[(df.Fare.notnull()) & (df.Survived == 0)], 'Fare', shade = True, color = 'red')
ax = sns.kdeplot(data = df.loc[(df.Fare.notnull()) & (df.Survived == 1)], 'Fare', shade = True, color = 'green')
ax.legend(['Not Survived', 'Survived'])
ax.set_xlim(0, 150)
```

Out[37]: (0.0, 150.0)



## 2.11: Cabin

Cabin has many missing values. We will keep first character from the cabin number and replace missing values with U as in Unknown.

```
In [38]: df.loc[df.Cabin.isnull(), 'Cabin'] = 'U'
```

```
In [39]: df['Cabin'] = df['Cabin'].map(lambda cab: cab[0:1])
```

```
In [40]: df.Cabin.value_counts()
```


```
Out[40]: U    1014
         C     94
         B     65
         D     46
         E     41
         A     22
         F     21
         G      5
         T      1
         Name: Cabin, dtype: int64
```

```
In [41]: # As G & T have very low numbers, we will classify those as Unknown as well
df['Cabin'].replace(['G', 'T'], 'U', inplace=True)
```

Cabin doesn't show much correlation with survival. Unknown cabin is showing low survival probability only because of high frequency.

```
In [42]: sns.catplot(x='Cabin', y='Survived', kind='bar', data=df, aspect=1.5)
```

```
Out[42]: <seaborn.axisgrid.FacetGrid at 0x1de95edfa08>
```



## Section 3: Feature Engineering

### 3.1: Encoding Categorical & Binary Columns

Before the encoding, we will drop the columns which are not required for modelling.

```
In [43]: df.drop(['PassengerId', 'Name'], axis=1, inplace=True)
df.shape
```

```
Out[43]: (1309, 11)
```

Let's identify the different type of columns first.



```
In [44]: numeric_cols = list(df.columns[df.dtypes == 'int64']) + list(df.columns[df.dtypes == 'float64'])
numeric_cols.remove('Survived')
binary_cols = list(['Sex'])
cat_cols = list(set(df.columns) - set(numeric_cols) - set(binary_cols))
cat_cols.remove('Survived')
```

```
In [45]: numeric_cols, binary_cols, cat_cols
```

```
Out[45]: (['Pclass', 'SibSp', 'Parch', 'Age', 'Fare'],
          ['Sex'],
          ['Ticket', 'Embarked', 'Salutation', 'Cabin'])
```

## Encoding Binary and Categorical Columns

```
In [46]: from sklearn.preprocessing import LabelBinarizer, MinMaxScaler

lb = LabelBinarizer()
for col in binary_cols:
    df[col] = lb.fit_transform(df[col])

df = pd.get_dummies(data = df, columns=cat_cols, drop_first=True)
```

Now, it's time to separate the labelled and unlabelled data, and create train-test splits.

```
In [47]: # Separate labelled and unlabelled data
unlabelled_df = df.loc[df.Survived.isna()]
X = df.loc[df.Survived.notna()].drop('Survived', axis=1)
y = df.loc[df.Survived.notna()]['Survived']
```

```
In [48]: # Create train-test splits
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state=0)
```

## 3.2: Scaling Numerical columns

```
In [50]: from sklearn.preprocessing import StandardScaler

sc = StandardScaler()

X_train_scaled = sc.fit_transform(X_train)
X_test_scaled = sc.transform(X_test)
```

# Section 4: Deep Learning Models

In this section, we will train multiple neural network models and compare performances using accuracy. Once we find the best model, we can check other metrics.

#### 4.1: Base model

We will generate Logistic Regression model as a base model. It will be used to compare the accuracies of the deep learning models.

```
In [51]: from sklearn.model_selection import GridSearchCV, KFold
from sklearn.linear_model import LogisticRegressionCV

kf = KFold(n_splits = 4, shuffle = True)

params = {'Cs': [2, 5, 10],
          'penalty': ['l1', 'l2'],
          'solver': ['newton-cg', 'lbfgs', 'liblinear']}

grid = GridSearchCV(estimator = LogisticRegressionCV(),
                    param_grid = params,
                    scoring = 'accuracy',
                    cv = kf,
                    n_jobs = -1)
grid.fit(X_train_scaled, y_train)
```

```
Out[51]: GridSearchCV(cv=KFold(n_splits=4, random_state=None, shuffle=True),
                      estimator=LogisticRegressionCV(), n_jobs=-1,
                      param_grid={'Cs': [2, 5, 10], 'penalty': ['l1', 'l2'],
                                   'solver': ['newton-cg', 'lbfgs', 'liblinear']}},
                      scoring='accuracy')
```

```
In [52]: print('Logistic Regression score: ', round(grid.best_score_, 4))
```

```
Logistic Regression score: 0.8329
```

#### 4.2: Basic Neural Network

A neural network with no hidden layer.

```
In [53]: from keras.models import Sequential
from keras.layers import Dense

nn1 = Sequential()
nn1.add(Dense(units=25, input_shape = (24,), activation='relu'))
nn1.add(Dense(units=1, activation='sigmoid'))
nn1.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 25)	625
dense_1 (Dense)	(None, 1)	26

Total params: 651  
 Trainable params: 651  
 Non-trainable params: 0

```
In [54]: nn1.compile(optimizer = 'rmsprop', loss = 'binary_crossentropy', metrics = 'accuracy')
nn1.fit(X_train_scaled, y_train, batch_size = 32, epochs = 20, validation_data
= (X_test_scaled, y_test))
```

Epoch 1/20  
23/23 [=====] - 0s 7ms/step - loss: 0.6622 - accuracy: 0.6067 - val\_loss: 0.5906 - val\_accuracy: 0.7486  
Epoch 2/20  
23/23 [=====] - 0s 2ms/step - loss: 0.5742 - accuracy: 0.7346 - val\_loss: 0.5229 - val\_accuracy: 0.7933  
Epoch 3/20  
23/23 [=====] - 0s 2ms/step - loss: 0.5225 - accuracy: 0.7697 - val\_loss: 0.4825 - val\_accuracy: 0.7933  
Epoch 4/20  
23/23 [=====] - 0s 2ms/step - loss: 0.4885 - accuracy: 0.7893 - val\_loss: 0.4520 - val\_accuracy: 0.7933  
Epoch 5/20  
23/23 [=====] - 0s 2ms/step - loss: 0.4671 - accuracy: 0.7963 - val\_loss: 0.4329 - val\_accuracy: 0.7933  
Epoch 6/20  
23/23 [=====] - 0s 2ms/step - loss: 0.4533 - accuracy: 0.7992 - val\_loss: 0.4197 - val\_accuracy: 0.7989  
Epoch 7/20  
23/23 [=====] - 0s 2ms/step - loss: 0.4438 - accuracy: 0.8048 - val\_loss: 0.4123 - val\_accuracy: 0.8101  
Epoch 8/20  
23/23 [=====] - 0s 2ms/step - loss: 0.4363 - accuracy: 0.8104 - val\_loss: 0.4054 - val\_accuracy: 0.8156  
Epoch 9/20  
23/23 [=====] - 0s 2ms/step - loss: 0.4304 - accuracy: 0.8202 - val\_loss: 0.4003 - val\_accuracy: 0.8156  
Epoch 10/20  
23/23 [=====] - 0s 2ms/step - loss: 0.4251 - accuracy: 0.8174 - val\_loss: 0.3973 - val\_accuracy: 0.8156  
Epoch 11/20  
23/23 [=====] - 0s 3ms/step - loss: 0.4203 - accuracy: 0.8272 - val\_loss: 0.3958 - val\_accuracy: 0.8212  
Epoch 12/20  
23/23 [=====] - 0s 2ms/step - loss: 0.4166 - accuracy: 0.8272 - val\_loss: 0.3932 - val\_accuracy: 0.8268  
Epoch 13/20  
23/23 [=====] - 0s 2ms/step - loss: 0.4127 - accuracy: 0.8301 - val\_loss: 0.3914 - val\_accuracy: 0.8268  
Epoch 14/20  
23/23 [=====] - 0s 2ms/step - loss: 0.4090 - accuracy: 0.8329 - val\_loss: 0.3898 - val\_accuracy: 0.8324  
Epoch 15/20  
23/23 [=====] - 0s 2ms/step - loss: 0.4056 - accuracy: 0.8371 - val\_loss: 0.3878 - val\_accuracy: 0.8324  
Epoch 16/20  
23/23 [=====] - 0s 2ms/step - loss: 0.4024 - accuracy: 0.8343 - val\_loss: 0.3869 - val\_accuracy: 0.8324  
Epoch 17/20  
23/23 [=====] - 0s 2ms/step - loss: 0.3999 - accuracy: 0.8385 - val\_loss: 0.3865 - val\_accuracy: 0.8380  
Epoch 18/20  
23/23 [=====] - 0s 2ms/step - loss: 0.3971 - accuracy: 0.8385 - val\_loss: 0.3850 - val\_accuracy: 0.8380  
Epoch 19/20  
23/23 [=====] - 0s 2ms/step - loss: 0.3947 - accuracy: 0.8385 - val\_loss: 0.3829 - val\_accuracy: 0.8492

```
Epoch 20/20
23/23 [=====] - 0s 1ms/step - loss: 0.3914 - accurac
y: 0.8385 - val_loss: 0.3816 - val_accuracy: 0.8492
```

```
Out[54]: <tensorflow.python.keras.callbacks.History at 0x1de9ccf4148>
```

We are able to match the accuracy just with a basic neural network. Let's try to improve on it.

### 4.3: Deep Neural Network

We will introduce one hidden layer.

```
In [55]: nn2 = Sequential()
nn2.add(Dense(units=25, input_shape = (24,), activation='relu'))
nn2.add(Dense(units=50, activation='relu'))
nn2.add(Dense(units=1, activation='sigmoid'))
nn2.summary()
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_2 (Dense)	(None, 25)	625
dense_3 (Dense)	(None, 50)	1300
dense_4 (Dense)	(None, 1)	51
=====	=====	=====
Total params: 1,976		
Trainable params: 1,976		
Non-trainable params: 0		
=====		

```
In [56]: nn2.compile(optimizer = 'rmsprop', loss = 'binary_crossentropy', metrics = 'accuracy')
nn2.fit(X_train_scaled, y_train, batch_size = 32, epochs = 20, validation_data
= (X_test_scaled, y_test))
```

Epoch 1/20  
23/23 [=====] - 0s 6ms/step - loss: 0.6345 - accuracy: 0.6334 - val\_loss: 0.5371 - val\_accuracy: 0.7709  
Epoch 2/20  
23/23 [=====] - 0s 2ms/step - loss: 0.5356 - accuracy: 0.7767 - val\_loss: 0.4699 - val\_accuracy: 0.8268  
Epoch 3/20  
23/23 [=====] - 0s 1ms/step - loss: 0.4867 - accuracy: 0.8006 - val\_loss: 0.4354 - val\_accuracy: 0.8436  
Epoch 4/20  
23/23 [=====] - 0s 2ms/step - loss: 0.4611 - accuracy: 0.8062 - val\_loss: 0.4188 - val\_accuracy: 0.8324  
Epoch 5/20  
23/23 [=====] - 0s 2ms/step - loss: 0.4431 - accuracy: 0.8146 - val\_loss: 0.4059 - val\_accuracy: 0.8380  
Epoch 6/20  
23/23 [=====] - 0s 2ms/step - loss: 0.4296 - accuracy: 0.8258 - val\_loss: 0.4002 - val\_accuracy: 0.8436  
Epoch 7/20  
23/23 [=====] - 0s 1ms/step - loss: 0.4185 - accuracy: 0.8258 - val\_loss: 0.3932 - val\_accuracy: 0.8436  
Epoch 8/20  
23/23 [=====] - 0s 1ms/step - loss: 0.4112 - accuracy: 0.8230 - val\_loss: 0.3881 - val\_accuracy: 0.8436  
Epoch 9/20  
23/23 [=====] - 0s 2ms/step - loss: 0.4037 - accuracy: 0.8244 - val\_loss: 0.3872 - val\_accuracy: 0.8547  
Epoch 10/20  
23/23 [=====] - 0s 2ms/step - loss: 0.3973 - accuracy: 0.8329 - val\_loss: 0.3845 - val\_accuracy: 0.8380  
Epoch 11/20  
23/23 [=====] - 0s 1ms/step - loss: 0.3921 - accuracy: 0.8329 - val\_loss: 0.3815 - val\_accuracy: 0.8380  
Epoch 12/20  
23/23 [=====] - 0s 2ms/step - loss: 0.3863 - accuracy: 0.8357 - val\_loss: 0.3810 - val\_accuracy: 0.8380  
Epoch 13/20  
23/23 [=====] - 0s 2ms/step - loss: 0.3829 - accuracy: 0.8399 - val\_loss: 0.3807 - val\_accuracy: 0.8380  
Epoch 14/20  
23/23 [=====] - 0s 2ms/step - loss: 0.3783 - accuracy: 0.8413 - val\_loss: 0.3783 - val\_accuracy: 0.8380  
Epoch 15/20  
23/23 [=====] - 0s 2ms/step - loss: 0.3747 - accuracy: 0.8469 - val\_loss: 0.3787 - val\_accuracy: 0.8380  
Epoch 16/20  
23/23 [=====] - 0s 2ms/step - loss: 0.3721 - accuracy: 0.8497 - val\_loss: 0.3740 - val\_accuracy: 0.8436  
Epoch 17/20  
23/23 [=====] - 0s 1ms/step - loss: 0.3690 - accuracy: 0.8497 - val\_loss: 0.3714 - val\_accuracy: 0.8492  
Epoch 18/20  
23/23 [=====] - 0s 2ms/step - loss: 0.3662 - accuracy: 0.8469 - val\_loss: 0.3735 - val\_accuracy: 0.8380  
Epoch 19/20  
23/23 [=====] - 0s 3ms/step - loss: 0.3626 - accuracy: 0.8539 - val\_loss: 0.3726 - val\_accuracy: 0.8436



```
Epoch 20/20  
23/23 [=====] - 0s 2ms/step - loss: 0.3613 - accurac  
y: 0.8455 - val_loss: 0.3733 - val_accuracy: 0.8380
```

```
Out[56]: <tensorflow.python.keras.callbacks.History at 0x1de9ef34ac8>
```

Let's see if accuracy improves by introducing more layers.


```
In [98]: nn2 = Sequential()
nn2.add(Dense(units=25, input_shape = (24,), activation='relu'))
nn2.add(Dense(units=50, activation='relu'))
nn2.add(Dense(units=50, activation='relu'))
nn2.add(Dense(units=1, activation='sigmoid'))
nn2.compile(optimizer = 'rmsprop', loss = 'binary_crossentropy', metrics = 'accuracy')
nn2_steps = nn2.fit(X_train_scaled, y_train, batch_size = 32, epochs = 20, validation_data = (X_test_scaled, y_test))
```

Epoch 1/20  
23/23 [=====] - 0s 7ms/step - loss: 0.6231 - accuracy: 0.6854 - val\_loss: 0.5368 - val\_accuracy: 0.7542  
Epoch 2/20  
23/23 [=====] - 0s 2ms/step - loss: 0.5203 - accuracy: 0.7781 - val\_loss: 0.4631 - val\_accuracy: 0.7933  
Epoch 3/20  
23/23 [=====] - 0s 3ms/step - loss: 0.4727 - accuracy: 0.7963 - val\_loss: 0.4351 - val\_accuracy: 0.8045  
Epoch 4/20  
23/23 [=====] - 0s 2ms/step - loss: 0.4492 - accuracy: 0.8048 - val\_loss: 0.4223 - val\_accuracy: 0.8101  
Epoch 5/20  
23/23 [=====] - 0s 2ms/step - loss: 0.4320 - accuracy: 0.8146 - val\_loss: 0.4154 - val\_accuracy: 0.8324  
Epoch 6/20  
23/23 [=====] - 0s 3ms/step - loss: 0.4181 - accuracy: 0.8244 - val\_loss: 0.4070 - val\_accuracy: 0.8324  
Epoch 7/20  
23/23 [=====] - 0s 3ms/step - loss: 0.4074 - accuracy: 0.8287 - val\_loss: 0.4112 - val\_accuracy: 0.8268  
Epoch 8/20  
23/23 [=====] - 0s 2ms/step - loss: 0.4004 - accuracy: 0.8413 - val\_loss: 0.4007 - val\_accuracy: 0.8380  
Epoch 9/20  
23/23 [=====] - 0s 2ms/step - loss: 0.3867 - accuracy: 0.8497 - val\_loss: 0.3935 - val\_accuracy: 0.8380  
Epoch 10/20  
23/23 [=====] - 0s 3ms/step - loss: 0.3799 - accuracy: 0.8455 - val\_loss: 0.3955 - val\_accuracy: 0.8436  
Epoch 11/20  
23/23 [=====] - 0s 2ms/step - loss: 0.3753 - accuracy: 0.8483 - val\_loss: 0.3892 - val\_accuracy: 0.8436  
Epoch 12/20  
23/23 [=====] - 0s 2ms/step - loss: 0.3677 - accuracy: 0.8497 - val\_loss: 0.3847 - val\_accuracy: 0.8603  
Epoch 13/20  
23/23 [=====] - 0s 2ms/step - loss: 0.3594 - accuracy: 0.8511 - val\_loss: 0.3823 - val\_accuracy: 0.8492  
Epoch 14/20  
23/23 [=====] - 0s 2ms/step - loss: 0.3580 - accuracy: 0.8511 - val\_loss: 0.3792 - val\_accuracy: 0.8492  
Epoch 15/20  
23/23 [=====] - 0s 2ms/step - loss: 0.3509 - accuracy: 0.8581 - val\_loss: 0.3728 - val\_accuracy: 0.8603  
Epoch 16/20  
23/23 [=====] - 0s 3ms/step - loss: 0.3463 - accuracy: 0.8610 - val\_loss: 0.3793 - val\_accuracy: 0.8603  
Epoch 17/20  
23/23 [=====] - 0s 2ms/step - loss: 0.3420 - accuracy: 0.8581 - val\_loss: 0.3836 - val\_accuracy: 0.8603  
Epoch 18/20  
23/23 [=====] - 0s 2ms/step - loss: 0.3394 - accuracy: 0.8581 - val\_loss: 0.3840 - val\_accuracy: 0.8380  
Epoch 19/20  
23/23 [=====] - 0s 2ms/step - loss: 0.3351 - accuracy: 0.8553 - val\_loss: 0.3806 - val\_accuracy: 0.8603

Epoch 20/20  
23/23 [=====] - 0s 2ms/step - loss: 0.3314 - accuracy: 0.8624 - val\_loss: 0.3712 - val\_accuracy: 0.8659

Let's plot the accuracies captured in each epochs.

```
In [110]: fig = plt.figure(figsize=(7, 5))
ax = fig.add_subplot(1,1,1)
ax.plot(nn2_steps.history['accuracy'], label = 'Training Accuracy')
ax.plot(nn2_steps.history['val_accuracy'], label = 'Test Accuracy')
plt.xticks(range(0,21,1))
plt.legend()
plt.tight_layout()
```



#### 4.4: Second Deep Neural Network

We will now try 'adam' optimizer, smaller batch size and more epochs.

```
In [105]: nn3 = Sequential()
nn3.add(Dense(units=25, input_shape = (24,), activation='relu'))
nn3.add(Dense(units=50, activation='relu'))
nn3.add(Dense(units=50, activation='relu'))
nn3.add(Dense(units=1, activation='sigmoid'))
nn3.compile(optimizer = 'adam', loss = 'binary_crossentropy', metrics = 'accuracy')
nn3_steps = nn3.fit(X_train_scaled, y_train, batch_size = 1, epochs = 50, validation_data = (X_test_scaled, y_test), shuffle=True)
```

Epoch 1/50  
712/712 [=====] - 1s 1ms/step - loss: 0.5127 - accuracy: 0.7697 - val\_loss: 0.4270 - val\_accuracy: 0.7877

Epoch 2/50  
712/712 [=====] - 1s 1ms/step - loss: 0.4391 - accuracy: 0.8132 - val\_loss: 0.4143 - val\_accuracy: 0.8101

Epoch 3/50  
712/712 [=====] - 1s 1ms/step - loss: 0.4138 - accuracy: 0.8272 - val\_loss: 0.3776 - val\_accuracy: 0.8380

Epoch 4/50  
712/712 [=====] - 1s 1ms/step - loss: 0.3906 - accuracy: 0.8371 - val\_loss: 0.4019 - val\_accuracy: 0.8492

Epoch 5/50  
712/712 [=====] - 1s 1ms/step - loss: 0.3743 - accuracy: 0.8511 - val\_loss: 0.4106 - val\_accuracy: 0.8045

Epoch 6/50  
712/712 [=====] - 1s 1ms/step - loss: 0.3700 - accuracy: 0.8413 - val\_loss: 0.3793 - val\_accuracy: 0.8324

Epoch 7/50  
712/712 [=====] - 1s 1ms/step - loss: 0.3559 - accuracy: 0.8581 - val\_loss: 0.3550 - val\_accuracy: 0.8659

Epoch 8/50  
712/712 [=====] - 1s 1ms/step - loss: 0.3529 - accuracy: 0.8539 - val\_loss: 0.3928 - val\_accuracy: 0.8436

Epoch 9/50  
712/712 [=====] - 1s 1ms/step - loss: 0.3377 - accuracy: 0.8624 - val\_loss: 0.3855 - val\_accuracy: 0.8547

Epoch 10/50  
712/712 [=====] - 1s 1ms/step - loss: 0.3445 - accuracy: 0.8553 - val\_loss: 0.3824 - val\_accuracy: 0.8380

Epoch 11/50  
712/712 [=====] - 1s 1ms/step - loss: 0.3274 - accuracy: 0.8722 - val\_loss: 0.4089 - val\_accuracy: 0.8492

Epoch 12/50  
712/712 [=====] - 1s 1ms/step - loss: 0.3274 - accuracy: 0.8624 - val\_loss: 0.4000 - val\_accuracy: 0.8380

Epoch 13/50  
712/712 [=====] - 1s 1ms/step - loss: 0.3208 - accuracy: 0.8638 - val\_loss: 0.3952 - val\_accuracy: 0.8268

Epoch 14/50  
712/712 [=====] - 1s 1ms/step - loss: 0.3196 - accuracy: 0.8638 - val\_loss: 0.3734 - val\_accuracy: 0.8268

Epoch 15/50  
712/712 [=====] - 1s 1ms/step - loss: 0.3057 - accuracy: 0.8708 - val\_loss: 0.4090 - val\_accuracy: 0.8324

Epoch 16/50  
712/712 [=====] - 1s 1ms/step - loss: 0.3146 - accuracy: 0.8708 - val\_loss: 0.3832 - val\_accuracy: 0.8212

Epoch 17/50  
712/712 [=====] - 1s 1ms/step - loss: 0.3092 - accuracy: 0.8694 - val\_loss: 0.4118 - val\_accuracy: 0.8212

Epoch 18/50  
712/712 [=====] - 1s 1ms/step - loss: 0.3002 - accuracy: 0.8736 - val\_loss: 0.4084 - val\_accuracy: 0.8324


Epoch 19/50  
712/712 [=====] - 1s 1ms/step - loss: 0.2985 - accuracy: 0.8778 - val\_loss: 0.4247 - val\_accuracy: 0.8212

Epoch 20/50  
712/712 [=====] - 1s 1ms/step - loss: 0.2952 - accuracy: 0.8820 - val\_loss: 0.4186 - val\_accuracy: 0.8212  
Epoch 21/50  
712/712 [=====] - 1s 1ms/step - loss: 0.2943 - accuracy: 0.8778 - val\_loss: 0.4245 - val\_accuracy: 0.8212  
Epoch 22/50  
712/712 [=====] - 1s 1ms/step - loss: 0.2910 - accuracy: 0.8764 - val\_loss: 0.3888 - val\_accuracy: 0.8380  
Epoch 23/50  
712/712 [=====] - 1s 1ms/step - loss: 0.2975 - accuracy: 0.8764 - val\_loss: 0.4564 - val\_accuracy: 0.8101  
Epoch 24/50  
712/712 [=====] - 1s 1ms/step - loss: 0.2794 - accuracy: 0.8848 - val\_loss: 0.4562 - val\_accuracy: 0.8268  
Epoch 25/50  
712/712 [=====] - 1s 2ms/step - loss: 0.2819 - accuracy: 0.8820 - val\_loss: 0.4440 - val\_accuracy: 0.8045  
Epoch 26/50  
712/712 [=====] - 1s 1ms/step - loss: 0.2860 - accuracy: 0.8862 - val\_loss: 0.4552 - val\_accuracy: 0.7989  
Epoch 27/50  
712/712 [=====] - 1s 2ms/step - loss: 0.2735 - accuracy: 0.8876 - val\_loss: 0.4640 - val\_accuracy: 0.8212  
Epoch 28/50  
712/712 [=====] - 1s 1ms/step - loss: 0.2709 - accuracy: 0.8904 - val\_loss: 0.4381 - val\_accuracy: 0.8268  
Epoch 29/50  
712/712 [=====] - 1s 1ms/step - loss: 0.2675 - accuracy: 0.8947 - val\_loss: 0.4499 - val\_accuracy: 0.8156  
Epoch 30/50  
712/712 [=====] - 1s 1ms/step - loss: 0.2710 - accuracy: 0.8834 - val\_loss: 0.4732 - val\_accuracy: 0.8045  
Epoch 31/50  
712/712 [=====] - 1s 1ms/step - loss: 0.2616 - accuracy: 0.8961 - val\_loss: 0.5229 - val\_accuracy: 0.7989  
Epoch 32/50  
712/712 [=====] - 1s 1ms/step - loss: 0.2630 - accuracy: 0.8961 - val\_loss: 0.5453 - val\_accuracy: 0.8045  
Epoch 33/50  
712/712 [=====] - 1s 1ms/step - loss: 0.2629 - accuracy: 0.8919 - val\_loss: 0.4945 - val\_accuracy: 0.8212  
Epoch 34/50  
712/712 [=====] - 1s 1ms/step - loss: 0.2568 - accuracy: 0.8947 - val\_loss: 0.4986 - val\_accuracy: 0.8212  
Epoch 35/50  
712/712 [=====] - 1s 2ms/step - loss: 0.2539 - accuracy: 0.8919 - val\_loss: 0.5162 - val\_accuracy: 0.8156  
Epoch 36/50  
712/712 [=====] - 1s 1ms/step - loss: 0.2493 - accuracy: 0.8933 - val\_loss: 0.5812 - val\_accuracy: 0.8101  
Epoch 37/50  
712/712 [=====] - 1s 1ms/step - loss: 0.2482 - accuracy: 0.8975 - val\_loss: 0.5447 - val\_accuracy: 0.8324  
Epoch 38/50  
712/712 [=====] - 1s 1ms/step - loss: 0.2517 - accuracy: 0.8975 - val\_loss: 0.5727 - val\_accuracy: 0.8212

Epoch 39/50  
712/712 [=====] - 1s 2ms/step - loss: 0.2482 - accuracy: 0.8947 - val\_loss: 0.5532 - val\_accuracy: 0.8156  
Epoch 40/50  
712/712 [=====] - 1s 1ms/step - loss: 0.2482 - accuracy: 0.9031 - val\_loss: 0.5444 - val\_accuracy: 0.8156  
Epoch 41/50  
712/712 [=====] - 1s 1ms/step - loss: 0.2428 - accuracy: 0.8933 - val\_loss: 0.5765 - val\_accuracy: 0.8156  
Epoch 42/50  
712/712 [=====] - 1s 1ms/step - loss: 0.2453 - accuracy: 0.8975 - val\_loss: 0.6245 - val\_accuracy: 0.8268  
Epoch 43/50  
712/712 [=====] - 1s 1ms/step - loss: 0.2412 - accuracy: 0.8989 - val\_loss: 0.5888 - val\_accuracy: 0.8268  
Epoch 44/50  
712/712 [=====] - 1s 1ms/step - loss: 0.2346 - accuracy: 0.8947 - val\_loss: 0.7620 - val\_accuracy: 0.8101  
Epoch 45/50  
712/712 [=====] - 1s 1ms/step - loss: 0.2939 - accuracy: 0.8876 - val\_loss: 0.5887 - val\_accuracy: 0.8212  
Epoch 46/50  
712/712 [=====] - 1s 1ms/step - loss: 0.2379 - accuracy: 0.8975 - val\_loss: 0.6382 - val\_accuracy: 0.8436  
Epoch 47/50  
712/712 [=====] - 1s 1ms/step - loss: 0.2248 - accuracy: 0.9059 - val\_loss: 0.6498 - val\_accuracy: 0.8101  
Epoch 48/50  
712/712 [=====] - 1s 1ms/step - loss: 0.2282 - accuracy: 0.9003 - val\_loss: 0.6508 - val\_accuracy: 0.8045  
Epoch 49/50  
712/712 [=====] - 1s 1ms/step - loss: 0.2250 - accuracy: 0.9017 - val\_loss: 0.6252 - val\_accuracy: 0.8101  
Epoch 50/50  
712/712 [=====] - 1s 1ms/step - loss: 0.2246 - accuracy: 0.9017 - val\_loss: 0.6986 - val\_accuracy: 0.8324



```
In [109]: fig = plt.figure(figsize=(7, 5))
ax = fig.add_subplot(1,1,1)
ax.plot(nn3_steps.history['accuracy'], label = 'Training Accuracy')
ax.plot(nn3_steps.history['val_accuracy'], label = 'Test Accuracy')
plt.xticks(range(0,51,5))
plt.legend()
plt.tight_layout()
```



#### 4.5: Model Selection & Performance

Our "Second Deep Neural Network" (nn3) is overfitted to the training data. So, we will stick to the first one (nn2) and measure other performance metrics.


```
In [130]: y_pred = nn2.predict(X_test_scaled)
y_pred[y_pred < 0.5] = 0
y_pred[y_pred >= 0.5] = 1
```

```
In [132]: from sklearn.metrics import confusion_matrix, classification_report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0.0	0.88	0.90	0.89	110
1.0	0.84	0.81	0.82	69
accuracy			0.87	179
macro avg	0.86	0.86	0.86	179
weighted avg	0.87	0.87	0.87	179

```
In [133]: cm = confusion_matrix(y_test, y_pred)
fig, ax = plt.subplots(figsize=(7, 5))
ax = sns.heatmap(cm, fmt='d', annot=True, cmap='Blues')
ax.set_xticklabels(['Not Survived', 'Survived'])
ax.set_yticklabels(['Not Survived', 'Survived'])
plt.xlabel('Predicted', fontsize=16)
plt.ylabel('Actual', fontsize=16)
```

Out[133]: Text(39.5, 0.5, 'Actual')



## Section 5: Predictions for Unlabelled data

As we don't know the labels, we will have to prepare and submit the prediction on Kaggle to get the accuracy score.

```
In [134]: X_unlabelled = unlabelled_df.drop('Survived', axis=1)
X_unlabelled_scaled = sc.transform(X_unlabelled)

y_unlabelled_pred = nn2.predict(X_unlabelled_scaled)
y_unlabelled_pred[y_unlabelled_pred < 0.5] = 0
y_unlabelled_pred[y_unlabelled_pred >= 0.5] = 1
```

```
In [147]: pid_df = pd.read_csv('test.csv')
pids = pid_df.iloc[:, 0]
results = pd.concat([pids.astype('int'), pd.Series(y_unlabelled_pred[:,0]).astype('int')], axis=1)
results.columns = ['PassengerId', 'Survived']
results.to_csv('titanic_result_7Dec.csv', index=False)
```

## Summary

We have noticed significant performance improvement with the neural network models. Even though, Age, Gender and Travel Class show very strong correlation with survival probability, I guess, there's always a luck factor.

***Note: I submitted the results in Kaggle for the unlabelled dataset and got accuracy of 0.75119.***

## Next steps:

What I haven't done in this notebook is to check correlation between features and multicollinearity. I also didn't check and fine-tune all the hyper-parameters. We may be able to improve the model performance by removing multicollinearity and having more precise values of model hyper-parameters.